

My Exam Prep:

- CNCF Kubernetes Class + labs
- K8S The Hard Way run through
- Run through all the tasks in the k8s docs
- Practice with systemd, journald, openssl, cfssl, and etcd
- Work through the sections in Walid's github list

Try the following exercises interactively:

Note - there are no answers here on purpose. You should be able to do these yourself using the minimal docs that you are allowed to use during the test. At a minimum this should train you on where to look for this info during the test, without notes.

1. Create a node that has a SSD and label it as such.
 - a. Create a pod that is only scheduled on SSD nodes.
2. Create 2 pod definitions: the second pod should be scheduled to run anywhere the first pod is running - 2nd pod runs alongside the first pod.
3. Create a deployment running nginx version 1.12.2 that will run in 2 pods
 - a. Scale this to 4 pods.
 - b. Scale it back to 2 pods.
 - c. Upgrade this to 1.13.8
 - d. Check the status of the upgrade
 - e. How do you do this in a way that you can see history of what happened?
 - f. Undo the upgrade
4. Create a service that uses a scratch disk.
 - a. Change the service to mount a disk from the host.
 - b. Change the service to mount a persistent volume.
5. Create a pod that has a liveness check
6. Create a service that manually requires endpoint creation - and create that too
7. Create a daemon set
 - a. Change the update strategy to do a rolling update but delaying 30 seconds between pod updates
8. Create a static pod
9. Create a busybox container without a manifest. Then edit the manifest.
10. Create a pod that uses secrets
 - a. Pull secrets from environment variables
 - b. Pull secrets from a volume
 - c. Dump the secrets out via kubectl to show it worked
11. Create a job that runs every 3 minutes and prints out the current time.
12. Create a job that runs 20 times, 5 containers at a time, and prints "Hello parallel world"
13. Create a service that uses an external load balancer and points to a 3 pod cluster running nginx.

14. Create a horizontal autoscaling group that starts with 2 pods and scales when CPU usage is over 50%.
15. Create a custom resource definition
 - a. Display it in the API with curl
16. Create a networking policy such that only pods with the label access=granted can talk to it.
 - a. Create an nginx pod and attach this policy to it.
 - b. Create a busybox pod and attempt to talk to nginx - should be blocked
 - c. Attach the label to busybox and try again - should be allowed
17. Create a service that references an externalname.
 - a. Test that this works from another pod
18. Create a pod that runs all processes as user 1000.
19. Create a namespace
 - a. Run a pod in the new namespace
 - b. Put memory limits on the namespace
 - c. Limit pods to 2 persistent volumes in this namespace
20. Write an ingress rule that redirects calls to /foo to one service and to /bar to another
21. Write a service that exposes nginx on a nodeport
 - a. Change it to use a cluster port
 - b. Scale the service
 - c. Change it to use an external IP
 - d. Change it to use a load balancer
22. Deploy nginx with 3 replicas and then expose a port
 - a. Use port forwarding to talk to a specific port
23. Make an API call using CURL and proper certs
24. Upgrade a cluster with kubeadm
25. Get logs for a pod
26. Deploy a pod with the wrong image name (like --image=nginy) and find the error message.
27. Get logs for kubectl
28. Get logs for the scheduler
29. Restart kubelet

Non-K8S

30. Convert a CRT to a PEM
 - a. Convert it back
31. Backup an etcd cluster
32. List the members of an etcd cluster
33. Find the health of etcd